

# Fetch Data Analyst Assessment

- Arjun Kumaran

## Part 1: Explore the data

Review the unstructured csv files and answer the following questions with code that supports your conclusions:

- Are there any data quality issues present?
- Are there any fields that are challenging to understand?

We recommend using SQL or python and data visualization to examine the data.

Data Sources:

1. PRODUCTS\_TAKEHOME.csv – Products Table
2. USER\_TAKEHOME.csv – Users Table
3. TRANSACTION\_TAKEHOME.csv – Transactions Table

Language used to analyse the data – **Python**

Firstly, a **basic data quality check** function was written commonly for all tables that does the following:

1. Counts null values and calculates percentage of null values in a column if null values are present in it.
2. Counts and drops duplicates if present.
3. Standardizes the text formatting by removing whitespaces and converting to lower-case.
4. Checks data-types.

### Products Table:

**Data Redundancy Issue:** There were 215 duplicates in the dataset and were removed.

### **Data Completeness Issue:**

```
Null values in products_df:
CATEGORY_1      111
CATEGORY_2     1424
CATEGORY_3     60566
CATEGORY_4     778093
MANUFACTURER    226474
BRAND           226472
BARCODE         4025
dtype: int64

Percentage of missing data in products_df:
CATEGORY_1      0.013128
CATEGORY_2      0.168411
CATEGORY_3      7.162895
CATEGORY_4     92.021898
MANUFACTURER    26.784160
BRAND           26.783923
BARCODE         0.476020
dtype: float64

Number of duplicate rows in products_df: 215
Duplicates dropped in products_df.
```

There is a lot of missing data in the dataset. For eg: 92% of data in the Category\_4 column is missing. After reviewing the data, I found that category\_1 is the main category. category\_2 is the subcategory of category 1. category\_3 is the subcategory of Category 2 and category 4 is the subcategory of category 3.

Missing data can break the links between categories and subcategories, making it difficult to understand the complete structure and relationships within the data.

Apart from having 226474 null values in the Manufacturer column, it has a placeholder string "placeholder manufacturer" 86895 times.

How should we handle these placeholder strings? Should they be considered as null values?

226472 values in the Brand column are also null values.

**Data Type Mismatches:** Numeric column "Barcode" is in object datatype. It had to be changed to INT.

## 2) Users Table:

### Data Completeness Issue:

```
Null values in user_df:
ID          0
CREATED_DATE 0
BIRTH_DATE   3675
STATE        4812
LANGUAGE     30508
GENDER       5892
dtype: int64

Percentage of missing data in user_df:
BIRTH_DATE    3.675
STATE         4.812
LANGUAGE     30.508
GENDER        5.892
dtype: float64
```

There is a lot of null values in this dataset. For eg: 30% of rows in Language are missing. This would impact the reliability of any user demographic analysis (Insights into their behavior, preferences, and trends based on language) done using this data.

**Data Type Mismatches:** Datetime columns "CREATED\_DATE" and "BIRTH\_DATE" are in object datatype. It had to be changed to datetime.

## Inconsistent Data:

user\_df.GENDER.value\_counts()

| GENDER                 | count |
|------------------------|-------|
| female                 | 64240 |
| male                   | 25829 |
| transgender            | 1772  |
| prefer_not_to_say      | 1350  |
| non_binary             | 473   |
| unknown                | 196   |
| not_listed             | 180   |
| non-binary             | 34    |
| not_specified          | 28    |
| my gender isn't listed | 5     |
| prefer not to say      | 1     |

There is inconsistency in the Gender terms and it had to be normalized.

```
[21] #Normalizing Gender terms
cleaned_user_df['GENDER'] = cleaned_user_df['GENDER'].str.replace('_', ' ')
cleaned_user_df['GENDER'] = cleaned_user_df['GENDER'].str.replace('-', ' ')
cleaned_user_df['GENDER'] = cleaned_user_df['GENDER'].str.replace('not Specified', 'not Listed')
cleaned_user_df['GENDER'] = cleaned_user_df['GENDER'].str.replace('my gender isn\'t listed', 'not listed')
```

cleaned\_user\_df['GENDER'].value\_counts() #Count after Gender Normalization

| GENDER            | count |
|-------------------|-------|
| female            | 64240 |
| male              | 25829 |
| transgender       | 1772  |
| prefer not to say | 1351  |
| non binary        | 507   |
| unknown           | 196   |
| not listed        | 185   |
| not specified     | 28    |

**Invalid Data (Range Violation):** To check invalid entries in the birth\_date column, age column was calculated using the following formula:

```
#Calculating Age of User
current_date = pd.Timestamp.now()
cleaned_user_df['age'] = (current_date - cleaned_user_df['BIRTH_DATE']).dt.days // 365
```

I assumed that if the age of the user was below 5 or above 100, it could be considered as invalid entries.

The Dataset had 75 invalid ages and they were removed from the dataset.

Following this check, the Age column was dropped from the dataset.

### 3) Transaction Table:

**Data Redundancy Issue:** There are 171 duplicate rows in this dataset and they have been removed.

#### Data Completeness Issue:

```
Null values in transaction_df:
RECEIPT_ID      0
PURCHASE_DATE   0
SCAN_DATE       0
STORE_NAME      0
USER_ID         0
BARCODE         5762
FINAL_QUANTITY  0
FINAL_SALE      0
dtype: int64

Percentage of missing data in transaction_df:
BARCODE    11.524
dtype: float64
```

There are null values only in the barcode column and it has been removed. Compared to the User and Products table, this table has very less null values. Upon Further Analysis, I found null values in the final quantity column.

#### Data Type Mismatches:

1. Barcode column had to be changed to INT Datatype.
2. Datetime columns "PURCHASE\_DATE" and "SCAN\_DATE" are in object datatype. It had to be changed to datetime.
3. Numeric columns "FINAL\_QUANTITY" and "FINAL\_SALE" are in object datatype. It had to be changed to Numeric.

#### Inconsistent Data:

```
cleaned_transaction_df.FINAL_QUANTITY.value_counts()
```

| FINAL_QUANTITY | count |
|----------------|-------|
| 1.00           | 31440 |
| zero           | 11067 |
| 2.00           | 1162  |
| 3.00           | 157   |
| 4.00           | 121   |
| ...            | ...   |
| 1.28           | 1     |
| 2.57           | 1     |
| 1.07           | 1     |
| 2.11           | 1     |
| 2.27           | 1     |

67 rows x 1 columns  
dtype: int64

In the FINAL\_QUANTITY column, It can be seen that 0 is represented as a string and it had to be changed as this is a numeric column.

```
[35] cleaned_transaction_df['FINAL_QUANTITY'] = cleaned_transaction_df['FINAL_QUANTITY'].replace('zero', 0.0) #Replacing "zero" to 0
```

cleaned\_transaction\_df['FINAL\_SALE'].value\_counts()

| FINAL_SALE | count |
|------------|-------|
| 11031      | 1     |
| 1.25       | 1106  |
| 1.00       | 675   |
| 2.00       | 510   |
| 2.99       | 500   |
| ...        | ...   |
| 8.64       | 1     |
| 7.45       | 1     |
| 10.79      | 1     |
| 13.14      | 1     |
| 7.83       | 1     |

1301 rows x 1 columns

dtype: int64

In the FINAL\_SALE column, It can be seen that null values are represented as a empty strings and it had to be changed as this is a numeric column.

```
[38] cleaned_transaction_df['FINAL_SALE'] = cleaned_transaction_df['FINAL_SALE'].replace('', np.nan) #Replacing empty string to null
```

**Date Constraint Violation:** Purchase Dates cannot be after Scan Dates as the user can scan receipts only after purchasing.

```
[42] #Sample Inconsistent dates
print(inconsistent_dates[['RECEIPT_ID', 'PURCHASE_DATE', 'SCAN_DATE']].head().to_string(index=False))
```

| RECEIPT_ID                           | PURCHASE_DATE | SCAN_DATE               |
|--------------------------------------|---------------|-------------------------|
| 008c1dcc-0f96-4b04-98c8-2a2bb63ef89d | 2024-07-21    | 2024-07-20 19:54:23.133 |
| 04a320ed-2903-45e5-8fd7-6eaf08daef32 | 2024-06-29    | 2024-06-28 11:03:31.783 |
| 05023b3d-5f83-47a7-a17c-8e8521d0bc94 | 2024-09-08    | 2024-09-07 22:22:29.903 |
| 06ce3da3-a588-4c37-93b4-0b6d11e42704 | 2024-06-22    | 2024-06-21 12:34:15.665 |
| 08d0e78f-3e63-40a3-8eb0-73fdf76da52c | 2024-06-22    | 2024-06-21 20:50:01.298 |

In the First row, it can be seen that the purchase date "2024-07-21" is after the scan date "2024-07-20". This is invalid because purchase date cannot be after scan date. So, this has to be removed.

The dataset had 47 inconsistent dates and had to be removed.

Data Integrity Issues:

The Transaction Table has User IDs that are not present in the User Table. This is logically not possible as a user has to create an account before scanning receipts through the app.

[ ] cleaned\_transaction\_df[cleaned\_transaction\_df['USER\_ID'] == '63b73a7f3d310dceeabd4758']

|  | RECEIPT_ID | PURCHASE_DATE                        | SCAN_DATE  | STORE_NAME              | USER_ID | BARCODE                  | FINAL_QUANTITY | FINAL_SALE |
|--|------------|--------------------------------------|------------|-------------------------|---------|--------------------------|----------------|------------|
|  | 41567      | 0000d256-4041-4a3e-adc4-5623fb6e0c99 | 2024-08-21 | 2024-08-21 14:19:06.539 | walmart | 63b73a7f3d310dceeabd4758 | 15300014978    | 1.0 1.54   |

[ ] cleaned\_user\_df[cleaned\_user\_df['ID'] == '63b73a7f3d310dceeabd4758']

|  | ID | CREATED_DATE | BIRTH_DATE | STATE | LANGUAGE | GENDER |
|--|----|--------------|------------|-------|----------|--------|
|--|----|--------------|------------|-------|----------|--------|

Upon reviewing the data, I found out that the transaction table has only only 85 unique user IDs in the transaction table that correspond to data in the User table. This is a big data integrity issue as the transaction table contains 15,879 Unique IDs that are not present in the User table. The dataset must be validated as this will impact any analysis made.

The Transaction Table also has Barcodes that are not present in the Products Table. While this is logically possible as some products might not be in Fetch’s product database, they should be added in the future.

[ ] cleaned\_products\_df[cleaned\_products\_df['BARCODE'] == 24000393429]

|  | CATEGORY_1 | CATEGORY_2 | CATEGORY_3 | CATEGORY_4 | MANUFACTURER | BRAND | BARCODE |
|--|------------|------------|------------|------------|--------------|-------|---------|
|--|------------|------------|------------|------------|--------------|-------|---------|

[ ] cleaned\_transaction\_df[cleaned\_transaction\_df['BARCODE'] == 24000393429]

|  | RECEIPT_ID | PURCHASE_DATE                        | SCAN_DATE  | STORE_NAME              | USER_ID | BARCODE                  | FINAL_QUANTITY | FINAL_SALE |
|--|------------|--------------------------------------|------------|-------------------------|---------|--------------------------|----------------|------------|
|  | 25004      | 21feab39-49f2-42e9-ae69-10371e2fc0a9 | 2024-08-23 | 2024-08-27 10:45:00.125 | target  | 61a58ac49c135b462ccdd1c  | 24000393429    | 1.0 2.59   |
|  | 36117      | 027f85a6-65aa-4dfd-a323-e10093953b66 | 2024-07-16 | 2024-07-16 15:31:40.795 | walmart | 5e718f029559aa138177814c | 24000393429    | 1.0 2.63   |
|  | 49658      | 05f0daef-b9ee-4e8b-9569-4387bce6886f | 2024-07-17 | 2024-07-31 13:07:46.658 | walmart | 5d6a7cd10d1b1f431660c97f | 24000393429    | 1.0 2.50   |

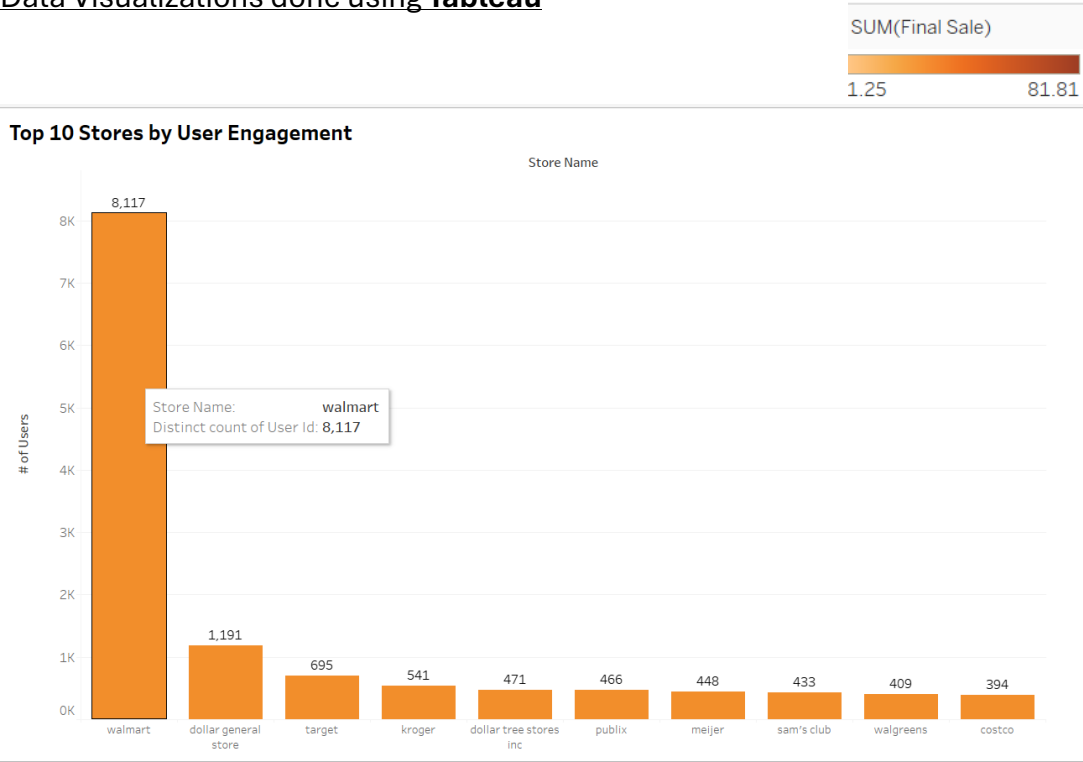
This dataframe has only 72 rows because it only has rows from the transaction table that have corresponding data in Users table(ID) and Products Table(Barcode).

product\_user\_transaction\_df

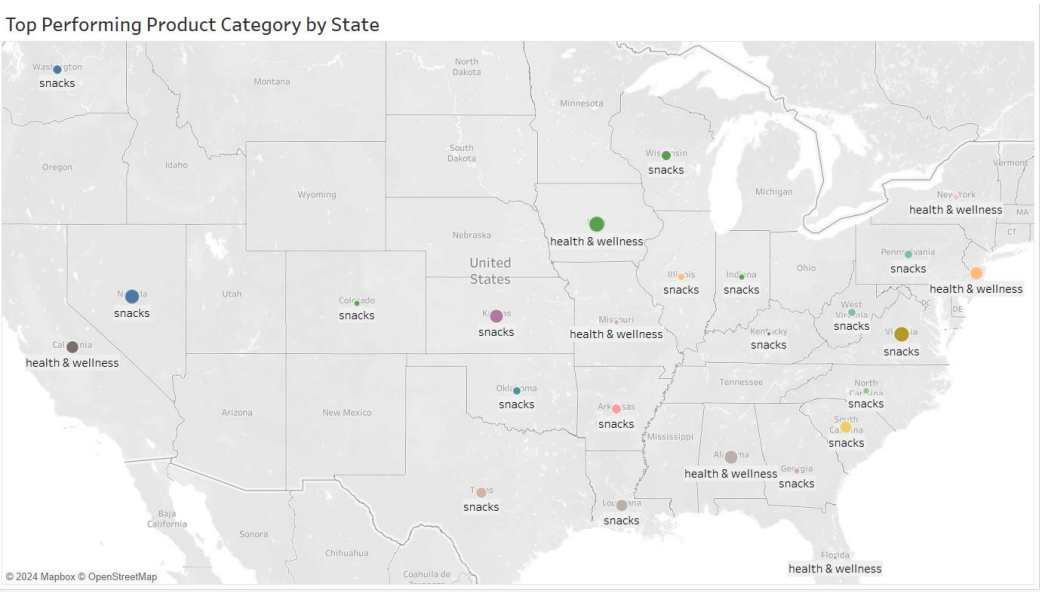
|       | RECEIPT_ID                           | PURCHASE_DATE | SCAN_DATE               | STORE_NAME             | USER_ID                  | BARCODE       | FINAL_QUANTITY | FINAL_SALE |
|-------|--------------------------------------|---------------|-------------------------|------------------------|--------------------------|---------------|----------------|------------|
| 25241 | 86fd365-1ddd-4f8b-bb6a-b924d4a54905  | 2024-08-03    | 2024-08-03 17:19:09.102 | walmart                | 63c8294d39c79dcbdd5c1e4f | 44000031114   | 1.0            | 3.88       |
| 25459 | c61751a9-9995-4eb3-8e39-d5fd7b08774f | 2024-08-11    | 2024-08-13 13:00:38.933 | save mart supermarkets | 5bb79c65b0a16836db35bbf1 | 311111423818  | 1.0            | 5.79       |
| 25489 | 9a6e4c08-b07b-4bad-ad44-330d2da1b67d | 2024-06-17    | 2024-06-17 11:20:43.843 | walmart                | 62815e99c907cf3f47d8ee35 | 810021200231  | 1.0            | 1.97       |
| 25712 | a713af20-8268-4094-8a41-3521e818bd15 | 2024-06-15    | 2024-06-15 12:54:12.582 | target                 | 61ca1f591a3acd27c1b18f0b | 28400517942   | 1.0            | 5.19       |
| 25874 | a7e6adf6-3dac-497f-90f1-7f371c639a1f | 2024-07-22    | 2024-07-22 09:49:41.406 | target                 | 5b441360be53340f289b0795 | 2700717433990 | 1.0            | 7.99       |
| ...   | ...                                  | ...           | ...                     | ...                    | ...                      | ...           | ...            | ...        |
| 47409 | 7297a22f-4a0a-4377-9eae-c8fb3e35ea4c | 2024-08-25    | 2024-08-25 15:42:48.533 | walmart                | 6615dab878ee6750bbc350ea | 817719021673  | 1.0            | 6.98       |
| 47670 | ad069462-7805-4c2b-a75f-283e5f63d851 | 2024-08-05    | 2024-08-13 07:53:18.199 | walmart                | 62c09104baa38d1a1f6c260e | 28400324434   | 1.0            | 2.32       |
| 48172 | d80f7f61-498f-46d7-be8a-0969971e2b7c | 2024-08-16    | 2024-08-17 08:55:14.932 | hy-vee                 | 5e6d0beb01ecdd13986852e8 | 30772112526   | 1.0            | 8.97       |
| 48691 | 47c724c5-680f-450d-a617-e36514ff06e  | 2024-07-18    | 2024-07-19 18:36:31.512 | h-e-b                  | 64ce5d823cb069b5eac9b700 | 41220078400   | 1.0            | 2.97       |
| 49556 | da8eda20-b96d-4df9-a618-45df75623eb6 | 2024-08-08    | 2024-08-11 06:43:27.935 | meijer                 | 65044dc5fe41d365c2ed7d71 | 760236261414  | 1.0            | 5.69       |

72 rows x 8 columns

Data Visualizations done using **Tableau**



From the above plot, it can be seen that Walmart is the store where most of Fetch’s active user base shops followed by Dollar General store. Increasing reward incentives at these top-performing stores will encourage continued loyalty and drive even more transactions. Exclusive promotions or bonus points at these stores could further boost user engagement.



From the above plot, it can be inferred that the top performing category (using sales) by state is Snacks followed by Health & Wellness. Targeted rewards campaigns for these categories could boost engagement.